

Project: Investigate a Dataset (TMDB Movie Data)

Table of Contents

- [Introduction](#)
- [Data Wrangling](#)
- [Exploratory Data Analysis](#)
- [Conclusions](#)

Introduction

This data set contains information about 10,000 movies collected from The Movie Database (TMDB), including user ratings and revenue. The main goal of this project is to explore the available data set and do the necessary wrangling and a reasonable conclusion. Below are some of the questions to be looked at;

- Does runtime affects the revenue?
- what kinds of properties are associated with movies that have high revenues(Top 20)
- Release year and revenue, is there a relationship between them.
- what year has the greatest revenue

```
In [2]: #Importing libraries

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [3]: import warnings
warnings.filterwarnings("ignore")
```

Data Wrangling

In this section of the report, I would be checking for cleanliness, and then trim and clean the dataset for analysis.

These process would include;

- Incorrect data types check
- Missing data check
- Removing duplicates from rows
- Renaming column names where necessary

Data Cleaning (General Properties)

```
In [4]: #Importing the dataset

df = pd.read_csv("tmdb-movies.csv")
```

```
In [5]: #Checking to know the available column names  
df.columns
```

```
Out[5]: Index(['id', 'imdb_id', 'popularity', 'budget', 'revenue', 'original_title',  
            'cast', 'homepage', 'director', 'tagline', 'keywords', 'overview',  
            'runtime', 'genres', 'production_companies', 'release_date',  
            'vote_count', 'vote_average', 'release_year', 'budget_adj',  
            'revenue_adj'],  
          dtype='object')
```

```
In [6]: #Investigating to know the information(non-null and dtype) of each column of the dataset  
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 10866 entries, 0 to 10865  
Data columns (total 21 columns):  
#   Column                                Non-Null Count  Dtype  
---  ---  
0   id                                    10866 non-null  int64  
1   imdb_id                             10856 non-null  object  
2   popularity                           10866 non-null  float64  
3   budget                               10866 non-null  int64  
4   revenue                              10866 non-null  int64  
5   original_title                       10866 non-null  object  
6   cast                                 10790 non-null  object  
7   homepage                             2936 non-null  object  
8   director                             10822 non-null  object  
9   tagline                              8042 non-null  object  
10  keywords                             9373 non-null  object  
11  overview                             10862 non-null  object  
12  runtime                              10866 non-null  int64  
13  genres                               10843 non-null  object  
14  production_companies                 9836 non-null  object  
15  release_date                         10866 non-null  object  
16  vote_count                           10866 non-null  int64  
17  vote_average                         10866 non-null  float64  
18  release_year                         10866 non-null  int64  
19  budget_adj                           10866 non-null  float64  
20  revenue_adj                           10866 non-null  float64  
dtypes: float64(4), int64(6), object(11)  
memory usage: 1.7+ MB
```

Fixing Datatypes: Changing Release Date datatype from object to date.

```
In [7]: def create_date_col(df, col):  
        """  
        This function helps to create a datetime column from an object  
        and returns datetime datatype  
        """  
        return pd.to_datetime(df[col])
```

```
In [8]: # convert release date to datetime using the column created above  
df['release_date'] = create_date_col(df, 'release_date')
```

```
In [9]: df.dtypes
```

```
Out[9]: id                                int64  
imdb_id                                object  
popularity                            float64  
budget                                int64  
revenue                               int64  
original_title                        object  
cast                                  object  
homepage                              object
```

```

director      object
tagline       object
keywords      object
overview      object
runtime       int64
genres        object
production_companies  object
release_date  datetime64[ns]
vote_count    int64
vote_average  float64
release_year  int64
budget_adj    float64
revenue_adj   float64
dtype: object

```

Number of null values in each column

```
In [10]: df.isnull().sum().sort_values(ascending=False)
```

```

Out[10]: homepage      7930
tagline      2824
keywords     1493
production_companies  1030
cast          76
director      44
genres        23
imdb_id       10
overview       4
budget_adj     0
release_year   0
vote_average   0
vote_count     0
release_date   0
id             0
runtime        0
original_title  0
revenue        0
budget         0
popularity     0
revenue_adj    0
dtype: int64

```

Filling null values with **no information** because forward fill, backward fill, mode and median fill is not reasonable for this type of data.

```
In [11]: df.fillna(value = "no information", inplace=True)
#confirming that there are no null values in the dataframe.
df.isnull().sum()
```

```

Out[11]: id      0
imdb_id  0
popularity  0
budget  0
revenue  0
original_title  0
cast  0
homepage  0
director  0
tagline  0
keywords  0
overview  0
runtime  0
genres  0
production_companies  0
release_date  0

```

```

vote_count      0
vote_average    0
release_year    0
budget_adj      0
revenue_adj     0
dtype: int64

```

```

In [12]: #Checking for duplicates in the dataframe
df.duplicated().sum()

```

```

Out[12]: 1

```

```

In [13]: #checking through the duplicated rows for confirmation before dropping.
df[df.duplicated(keep=False)]

```

```

Out[13]:

```

	id	imdb_id	popularity	budget	revenue	original_title	cast	homepage	director
2089	42194	tt0411951	0.59643	30000000	967000	TEKKEN	Jon Foo Kelly Overton Cary-Hiroyuki Tagawa lan...	no information	Dwight H. Little
2090	42194	tt0411951	0.59643	30000000	967000	TEKKEN	Jon Foo Kelly Overton Cary-Hiroyuki Tagawa lan...	no information	Dwight H. Little

2 rows x 21 columns

```

In [14]: #Dropping the duplicates after confirmation
df.drop_duplicates(inplace=True)

#checking to see if there are more duplicates after dropping
df.duplicated().sum()

```

```

Out[14]: 0

```

Exploratory Data Analysis

Research Question 1 (How does the runtime affects the revenue?)

```


In [15]: # define a function for all your plots
def axis_title_plot(x_label, y_label, title):

    """generate x label, y label and title for a plot
    it takes in the x_label, y_label, title
    and print out its attributes on a plot
    """

    plt.xlabel(x_label)
    plt.ylabel(y_label)
    plt.title(title)
    plt.show()

```

```
In [16]: # plot the distribution of revenue to understand how they align
df['revenue'].hist()
axis_title_plot(x_label='Revenue', y_label='Count', title='Distribution of Movie Revenue')
```




Summary of Findings: The distribution of the revenue from the graph is right-skewed, where most of the values (Revenue) are less than 100M. Meaning that only few movies can get a Revenue greater than 100M.

```
In [17]: df['revenue'].mean()
```

```
Out[17]: 39826896.07685228
```

```
In [18]: # Check the relationship between revenue and runtime using scatter plot
df.plot.scatter(x = 'revenue', y = 'runtime', figsize=(8,6))
axis_title_plot(x_label='Revenue', y_label='Runtime', title='Scatter Plot of Revenue vs
```



Summary of Findings: From the diagram above, it can be deduced that increase in runtime does affect the an increase in revenue. However we can see that majority of the Runtime id between the minute of 100 - 200 . We can also see big money Movies at runtime less than 200 Minutes.

Research Question 2 (Finding the relationship between the release year and revenue)

```
In [19]: #Looking at the top 10 revenues and their year of release using groupby Value  
df.groupby(['release_year']).mean()['revenue'].sort_values(ascending=False).head(10)
```

```
Out[19]: release_year  
1997      5.549569e+07  
2001      5.541357e+07  
2002      5.505120e+07  
2004      5.470301e+07  
2003      5.387275e+07  
1995      5.232195e+07  
1999      5.069515e+07  
2000      4.836432e+07  
1992      4.570040e+07  
2005      4.537592e+07  
Name: revenue, dtype: float64
```

Summary of Findings: Looking at the data above, late 90's and early 2000's have the highest revenue with 1997 leading the race.

Research Question 3 (What kinds of properties are associated with movies that have high revenue?)

```
In [20]: #I am trying to get all the rows associated with the 20 highest revenues  
df_high = df.sort_values(by=['revenue'], ascending=False).head(20)
```

```
In [21]: df_high.mean()
```

```
Out[21]: id                5.712840e+04  
popularity    7.184929e+00  
budget        1.883500e+08  
revenue        1.366390e+09  
runtime        1.382500e+02  
vote_count     4.515250e+03  
vote_average   6.990000e+00  
release_year   2.010000e+03  
budget_adj     1.868408e+08  
revenue_adj    1.387102e+09  
dtype: float64
```

Summary of Findings: Looking at the characteristics of the top 20 variables, we can deduce that the top revenues have an average runtime of 138 minutes and budget of 188million dollars.

Research Question 4 (What year has the greatest revenue?)

```
In [22]: # get movie with highest average revenue  
df.groupby(['release_year']).mean()['revenue'].sort_values(ascending=False).head(1)
```

```
Out[22]: release_year  
1997      5.549569e+07  
Name: revenue, dtype: float64
```

Summary of Findings: The year 1997 had the highest average revenue compared to other years. From the data, there had been big boom in movie sales for the Movie industry during this period.

Conclusions

Limitation

- I filled in missing values with 'no information', there could have been a better way to fill this.
- Most of my conclusions were estimated values because the values in the data set are too big.

Summary From Question

- **How does the runtime affects the revenue?**
 - The distribution of the revenue from the graph is right-skewed, where most of the values(Revenue) are less than 100M. *Meaning that only few movie can get a Revenue greater than '100M'*
- **Finding the relationship between the release year and revenue.**
 - From the plotted , it can be deduced that increase in runtime does affect the an increase in revenue. However we can see that majority of the Runtime id between the minute of 100 - 200 . We can also see big money Movies at runtime less than 200 Minutes. And also, Looking at the data above, late 90's and early 2000's have the highest revenue with 1997 leading the race.
- **What kinds of properties are associated with movies that have high revenue?**
 - Looking at the characteristics of the top 20 variables, we can deduce that the top revenues have an average runtime of 138 minutes and Revenue of 188 million dollars .
- **What year has the greatest revenue?**
 - The year 1997 had the highest average revenue compared to other years. From the data, there had been big boom in movie sales for the Movie industry during this period

Final Conclusion:

Looking at the exploration done in this dataset, one would find out that there were high revenues back in the early 90's and late 20's years compared to recent times, which means people watch movies more in the past than now that brought about more returns for the industry.